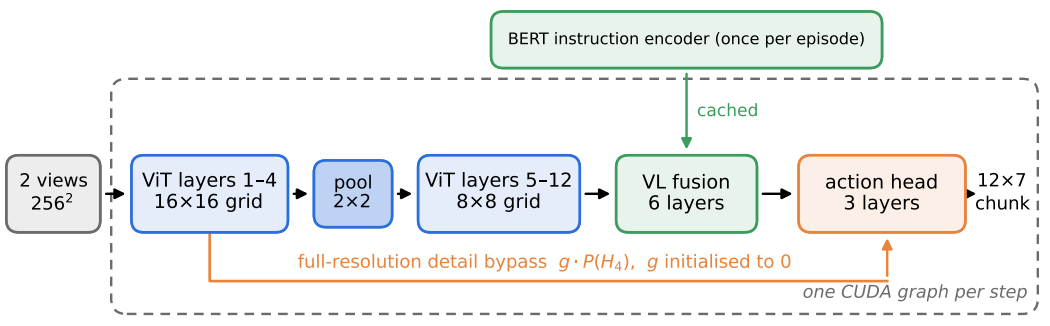
BERT instruction encoder (once per episode)
cached
2 views
256²
ViT layers 1–4
16×16 grid
pool
2×2
ViT layers 5–12
8×8 grid
VL fusion
6 layers
action head
3 layers
12×7
chunk
full-resolution detail bypass $g \cdot P(H_4)$, g initialised to 0
one CUDA graph per step